Stage A: prepare data once

2,000 questions → **3 teachers** × 3 samples → **18,000** trajectories → **answer** filtering → **Qual(t)** judge + rules

correct candidates (the student reads every one)

cached Qual(t), reused by every method

Stage B: per student × per method

Fam(t, m) once per student → **select k = 3** one rule per method → **5,685 samples** same size for all → **LoRA** fine-tuning

next method: reselect and retrain

next student: recompute Fam, then repeat every method

fine-tuned LoRA adapter, one per student × method × seed

Stage C: evaluate and deploy

6 benchmarks Acc@4, Pass@4 → **GGUF** Ollama, LM Studio