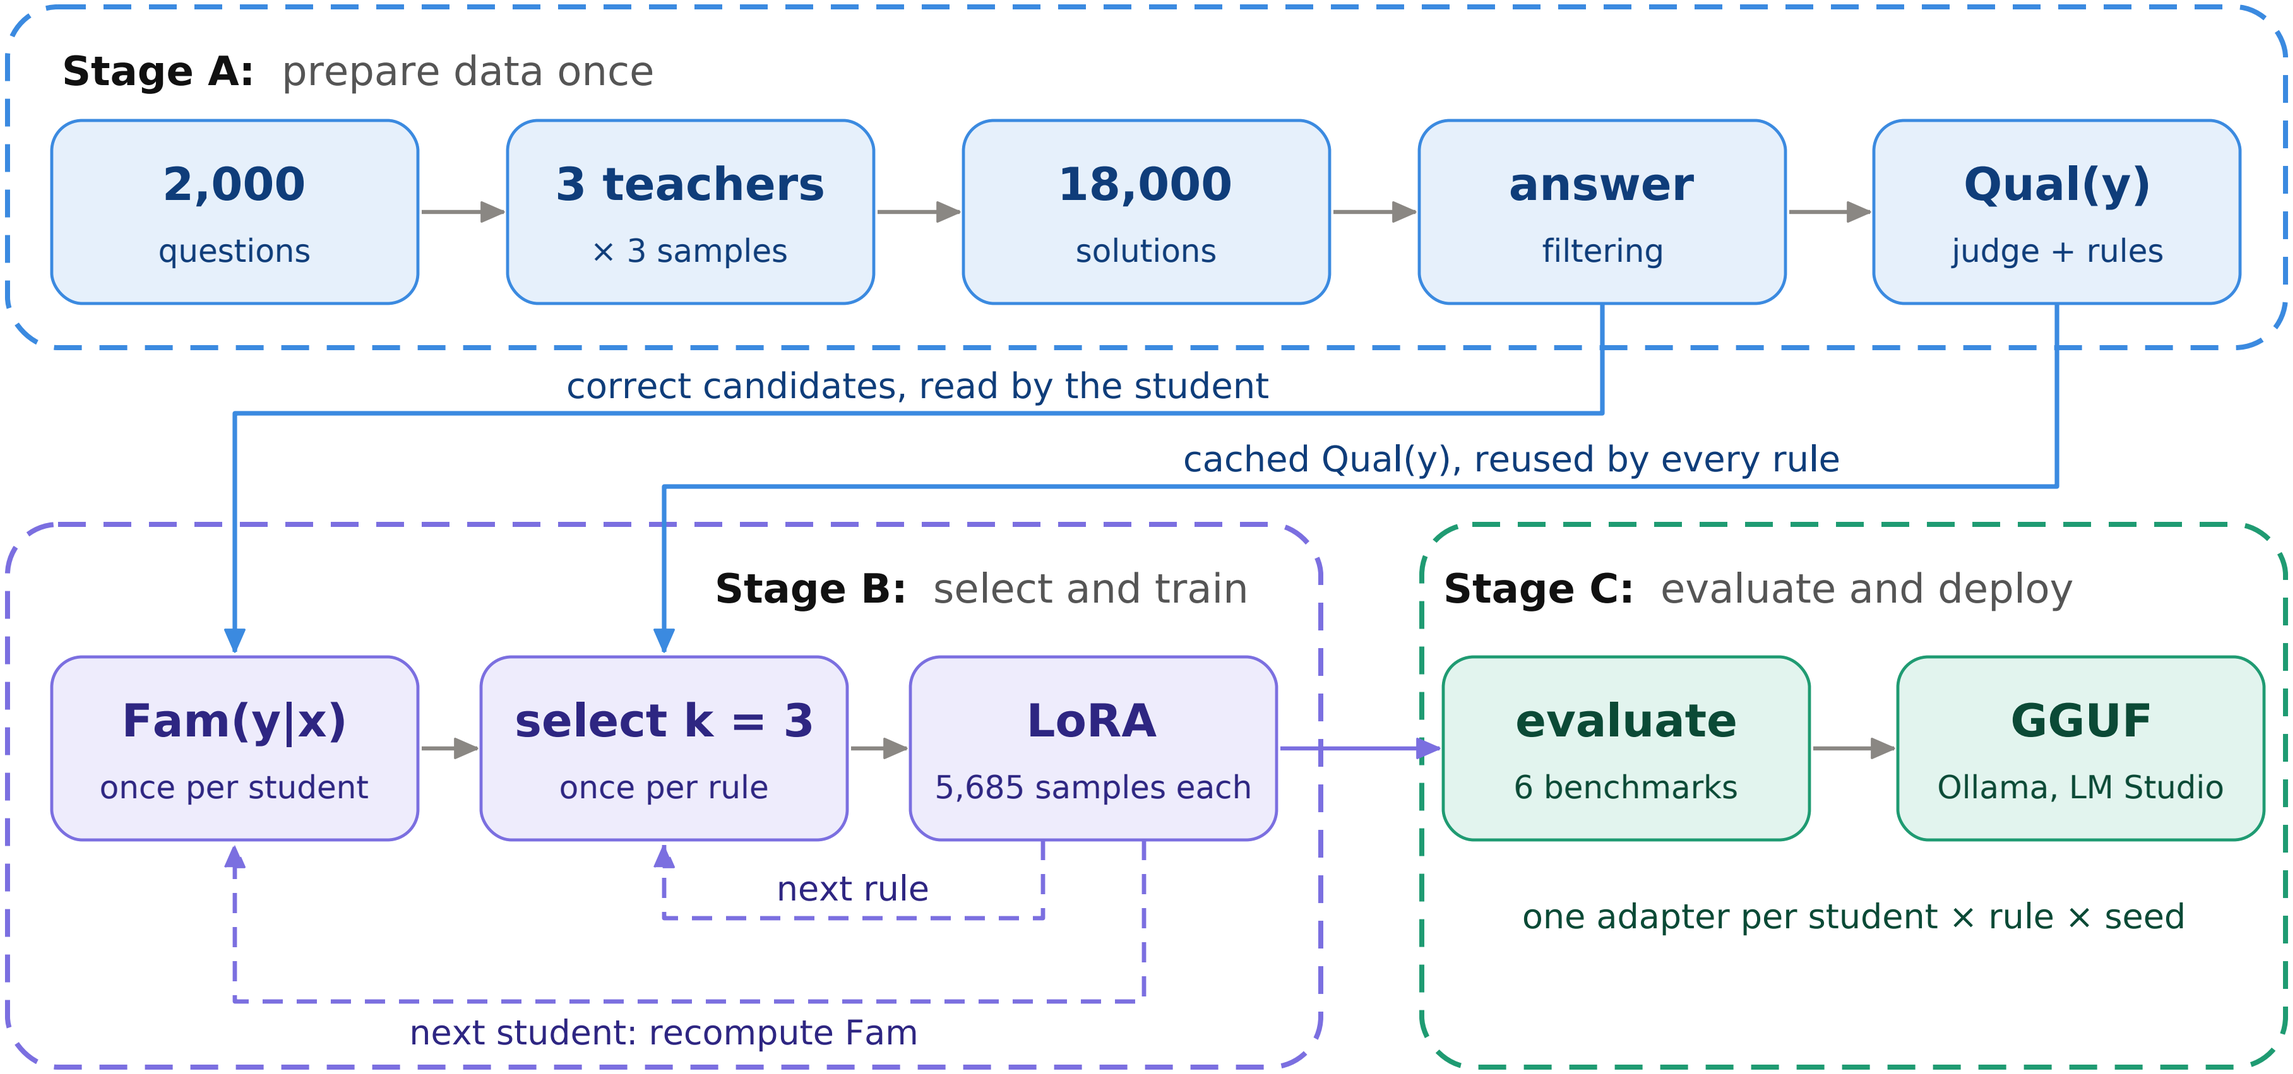
Stage A: prepare data once
2,000
questions
3 teachers
× 3 samples
18,000
solutions
answer
filtering
Qual(y)
judge + rules
correct candidates, read by the student
cached Qual(y), reused by every rule
Stage B: select and train
Fam(y|x)
once per student
select k = 3
once per rule
LoRA
5,685 samples each
next rule
next student: recompute Fam
Stage C: evaluate and deploy
evaluate
6 benchmarks
GGUF
Ollama, LM Studio
one adapter per student × rule × seed